

Basic results on λ -calculus

Review: λ -calculus syntax

- ▶ Let's review the basic notions of λ -calculus from last time.
- ▶ If V is a set of variables the set of valid λ -terms Λ is described by the grammar

$$\Lambda = V \mid \Lambda\Lambda \mid (\lambda V.\Lambda)$$

- ▶ We think of

$$\lambda x.U$$

as a function that takes x to U (where x can appear in U).

Currying

- ▶ As a shorthand, we write

$$\lambda x.(\lambda y.U) \quad \text{as} \quad \lambda xy. U$$

- ▶ The idea (known as “currying”) is that a function that returns a function can be thought of as a function of two variables.

Bound and free variables

- ▶ Any variable that appears next to the λ sign is called a **bound** variable

(technically the λ introduces a **binder**, and it is the occurrences in scope that are **bound**, but we lose nothing with this abuse of language)

- ▶ All other variables are called **free** variables.
- ▶ Thus in

$$\lambda xy. xyz$$

the variables x and y are bound, and z is free.

Bound variables in mathematics

- ▶ We can think of free variables as constants outside of our control.
- ▶ We regularly encounter bound variables in mathematics:

$$\int_0^1 f(x) dx \quad \lim_{x \rightarrow 0} f(x) \quad \sum_{i=1}^n a_i$$

- ▶ In each expression, the variable (x or i) is a placeholder — renaming it doesn't change the meaning.

α -conversion

- ▶ Renaming bound variables does not change the meaning of the λ -term:

$$\lambda x. xy = \lambda u. uy$$

- ▶ One can reuse bound variable names:

$$\lambda x. (\lambda x. x) = \lambda u. (\lambda x. x)$$

Both represent a function that returns the identity.

- ▶ We identify λ -terms that differ only by renaming bound variables. This process is called **α -conversion**.

In theoretical treatments, we often assume all bound variables have distinct names (the **Barendregt convention**) to simplify proofs about substitution. In implementations, techniques like de Bruijn indices or unique identifier generation serve a similar purpose.

β -reduction

- ▶ The computational rule of λ -calculus is **β -reduction**, which takes any expression of the form

$$(\lambda x.N) M$$

and evaluates it to

$$N[x := M]$$

i.e., the λ -term with every *free* instance of x in N replaced by M .

- ▶ **IMPORTANT** : If you have something like A B C the computation is always left-associative, i.e

$$ABC = (AB)C$$

$$ABCD = ((AB)C)D$$

Encoding computation

- ▶ λ -calculus consists of a syntax (valid expressions) and a single evaluation rule.
- ▶ To use it as a computational model, we first encode what we want to express:

$$\mathbf{true} = \lambda xy. x$$


$$\mathbf{false} = \lambda xy. y$$

$$\mathbf{if}\ b\ \mathbf{then}\ t\ \mathbf{else}\ e = b\ t\ e$$

- ▶ The conditional simply applies b to both branches — the boolean selects which one to return.

Verifying the encoding

We must verify our encoding behaves as intended:



Similarly: $\mathbf{false} \ t \ e = (\lambda xy. y) t e \xrightarrow{\beta} (\lambda y. y) e \xrightarrow{\beta} e \checkmark$

Semantics vs syntax


- ▶ We see here a fundamental concern of logic: the consistency between **semantics** (intended meaning) and **syntax** (formal rules). This connection will resurface when we study the relationship between λ -calculus and logic.
- ▶ For our boolean encoding, it's easy to verify the match.
- ▶ But for more complex systems, there might be a mismatch we fail to see — this would render our computations meaningless.

The consistency problem

- ▶ A basic worry: perhaps λ -calculus is itself inconsistent, and evaluating in different orders gives different results.
- ▶ If so, it wouldn't matter what our conventions are — the entire system would be useless!
- ▶ And this worry is justified...

An important example

Let $\Omega := (\lambda x. x x)(\lambda x. x x)$. Note that reducing Ω yields Ω again.



Different reduction orders give different behaviors! Can we still trust the results?

Two key questions

This example raises two fundamental questions:

1. **When does computation conclude?**
 - ▶ We need a notion of “final answer”
2. **Is the final answer unique?**
 - ▶ If different reduction paths could yield different final answers, the calculus would be useless

Today we'll answer both questions.

Toward a notion of “final answer”

- ▶ We will now define a **normal form** for λ -terms — a state that can reasonably be considered the final result of a computation.
- ▶ For this definition we need the notion of a **subterm**.
- ▶ For this purpose it is useful to realize that λ -terms can be decomposed as trees.

λ -terms as trees


Every λ -term can be represented as a tree:

- ▶ **Variables** x are leaves
- ▶ **Applications** $M N$ are binary nodes with M and N as children
- ▶ **Abstractions** $\lambda x.M$ are unary nodes with M as child

A **subterm** is any term corresponding to a subtree.

Tree decomposition: example


The term $(\lambda x.x y) z$ as a tree:



This is a **redex**: the root is an application whose left child is a λ -abstraction, i.e., a redex is any λ -term of the form $(\lambda x.N) M$ (a *reducible expression*).

Subterms

The **subterms** of $(\lambda x.x y) z$ are all the subtrees:




6 subterms total:

1. $(\lambda x.x y) z$
2. $\lambda x.x y$
3. z
4. $x y$
5. x
6. y

Finding redexes in a term

A term is in **β -normal form** if no subterm is a redex.



The term $(\lambda f.f x)((\lambda y.y)z)$ has **two redexes** and is **not** in β -normal form.

We think of β -normal forms as the final outcome of a computation — once we reach this stage, the computation must stop.

Examples

- The following are in β -normal form

$$xy, x(\lambda u.u)z$$

The second example is in normal form because it is interpreted as

$$(x(\lambda u.u))z$$

and $x(\lambda u.u)$ has no further reduction.

- The following is not in β -normal form

$$x((\lambda u.u)z), (\lambda u.xu)v$$

The second example is obvious, but the first is not in β -normal form because we have indicated with parenthesis that we should view $x((\lambda u.u)z)$ as MN with $M = x$ and $N = (\lambda u.u)z$ and N is a redex that further reduces to z .

Can we always reach a normal form?

- ▶ We now have a notion of “final answer” — the β -normal form.
- ▶ But can we always reach it? And if so, is it unique?
- ▶ To address these questions precisely, we need some notation.

Notation

- ▶ If N is obtained from M by a **single** β -reduction step:

$$M \rightarrow_{\beta} N$$

- ▶ If N can be obtained from M by a (finite) **chain** of β -reductions:

$$M \rightarrow_{\beta} M_1 \rightarrow_{\beta} M_2 \rightarrow_{\beta} \cdots \rightarrow_{\beta} N$$

then we write $M \twoheadrightarrow_{\beta} N$ (multi-step reduction).

- ▶ We think of β -normal forms as final answers — once reached, computation must stop.


Normalizable terms

- ▶ A λ -term M is **weakly normalizable** if there exists some reduction path to a normal form.
- ▶ A λ -term M is **strongly normalizable** if every reduction path terminates (no infinite paths exist).
- ▶ Strongly normalizable \Rightarrow weakly normalizable (but not conversely).

Weak vs strong normalization

Weakly Normalizable


\exists a path to normal form




At least one path terminates.

Strongly Normalizable

All paths terminate



No infinite reduction paths.



Strongly normalizable \subsetneq Weakly normalizable

β -equivalence

- ▶ Calculations don't have to go forward only.
- ▶ Two terms M and N are **β -equivalent**, written $M =_{\beta} N$, if there is a chain where each step is either \rightarrow_{β} or \leftarrow_{β} :

$$M := M_0 \rightarrow_{\beta} M_1 \leftarrow_{\beta} M_2 \rightarrow_{\beta} M_3 \rightarrow_{\beta} M_4 := N$$

- ▶ A λ -term M has a **β -normal form** if $M =_{\beta} N$ for some N in β -normal form.

Today's main results

Theorem (Main Theorem)

1. If $M =_{\beta} N$ for some N in β -normal form, then $M \rightarrow_{\beta} N$.
2. For any λ -term M , there is at most one β -normal form.

Significance:

- ▶ Normal forms can always be reached by forward calculation
- ▶ The result is unique (if it exists)
- ▶ If no normal form exists, the computation loops forever

The Church-Rosser theorem


Our main tool is the **Church-Rosser theorem** (also called **confluence**).

Theorem (Church-Rosser)

If $M \twoheadrightarrow_{\beta} N_1$ and $M \twoheadrightarrow_{\beta} N_2$, then there exists P such that $N_1 \twoheadrightarrow_{\beta} P$ and $N_2 \twoheadrightarrow_{\beta} P$.

Intuition: Reduction is like water flowing downhill — different paths may diverge, but they can always rejoin. No matter which redexes you choose to reduce, you can never paint yourself into a corner.

Church-Rosser: the diamond property



We will prove Church-Rosser next time. For now, let's use it.

Result 1: Uniqueness of normal forms

Theorem


If $M \twoheadrightarrow_{\beta} N_1$ and $M \twoheadrightarrow_{\beta} N_2$ with N_1, N_2 in β -normal form, then $N_1 = N_2$.

Proof: By Church-Rosser, $\exists L$ with $N_1 \twoheadrightarrow_{\beta} L$ and $N_2 \twoheadrightarrow_{\beta} L$.

But N_1 is in normal form, so $N_1 = L$. Similarly $N_2 = L$.

Therefore $N_1 = N_2$. \square

Uniqueness: diagram



$$N_1, N_2 \text{ normal} \Rightarrow N_1 = L = N_2$$

Result 2: Forward calculation

Theorem

If $M =_{\beta} N$ and N is in β -normal form, then $M \rightarrow_{\beta} N$.


To prove this, we need an auxiliary lemma:

Lemma (Auxiliary)

If $M =_{\beta} N$, then $\exists L$ such that $M \rightarrow_{\beta} L$ and $N \rightarrow_{\beta} L$.

Proof of theorem (assuming lemma): Given $M =_{\beta} N$ with N normal, the lemma gives us L with $M \rightarrow_{\beta} L$ and $N \rightarrow_{\beta} L$. Since N is normal, $N = L$. Thus $M \rightarrow_{\beta} N$. \square

Summary: from equivalence to reduction




Lemma gives L . But N normal $\Rightarrow N = L \Rightarrow M \rightarrow\!\!\!\rightarrow \beta N$

Proving the auxiliary lemma

Lemma

If $M =_{\beta} N$, then $\exists L$ such that $M \rightarrowtail_{\beta} L$ and $N \rightarrowtail_{\beta} L$.



Proof: By induction on the length of the $=_{\beta}$ chain.


The conversion chain

$$M \rightarrow M_1 \leftarrow M_2 \cdots M_k \rightarrow N$$

- ▶ Each arrow is either \rightarrow_β or ${}_\beta\leftarrow$
- ▶ Chain length = $k + 1$ (number of steps)

Base Case: Chain Length = 0

$M = N$ (identical terms)



$$L = M = N$$

(0 reduction steps each)

Take $L = M = N$. Both reduce to L in zero steps. ✓

Inductive Step: Chain Length = n

Assume true for chains of length $< n$.

The first step is either:

Case 1:

$$M \rightarrow_{\beta} P$$

(forward step)


Case 2:

$$M \xleftarrow{\beta} P$$

(backward step)


In both cases, $P =_{\beta} N$ with a shorter chain (length $n - 1$).

Case 1: $M \rightarrow_{\beta} P$ — Setup




The chain $P =_{\beta} N$ has length $n - 1$.

Case 1: $M \rightarrow_{\beta} P$ — Apply IH




Case 1: $M \rightarrow_{\beta} P$ — Conclusion




Since $M \rightarrow_{\beta} P \twoheadrightarrow_{\beta} L$, we have $M \twoheadrightarrow_{\beta} L$ and $N \twoheadrightarrow_{\beta} L$. ✓

Case 2: $M \xleftarrow{\beta} P$ — Setup




The chain $P =_\beta N$ has length $n - 1$.

Case 2: $M_\beta \leftarrow P$ — Apply IH




Case 2: $M \xrightarrow{\beta} P$ — Apply Church-Rosser



Apply Church-Rosser to $P \xrightarrow{\beta} M$ and $P \xrightarrow{\beta} L'$.


Case 2: $M \xleftarrow{\beta} P$ — Conclusion



$M \twoheadrightarrow_{\beta} L$ by CR. $N \twoheadrightarrow_{\beta} L'$ $\twoheadrightarrow_{\beta} L$. ✓

Proof complete

By induction, for all $M =_{\beta} N$, there exists L with $M \rightarrowtail_{\beta} L$ and $N \rightarrowtail_{\beta} L$.



□

Summary of results


Theorem (Main Results)

1. **Uniqueness:** Any λ -term has at most one β -normal form.
2. **Forward calculation:** If $M =_{\beta} N$ with N normal, then $M \rightarrow_{\beta} N$.

Consequences:

- ▶ β -normal forms are well-defined “final answers”
- ▶ We can always compute forward to reach them
- ▶ Different reduction strategies may take different paths, but the destination is unique

Logical dependencies



Next time

Proving the Church-Rosser theorem

The proof uses a clever technique:

1. Define **parallel reduction** \Rightarrow_{β} (reduce multiple redexes simultaneously)
2. Show parallel reduction satisfies the **diamond property** directly
3. Show \rightarrow_{β} is the transitive closure of \Rightarrow_{β}
4. Conclude Church-Rosser for \rightarrow_{β}

This is one of the most elegant proofs in the theory of computation!